

Daily Interest Digest — 2026-10-09 (FORMAT EXAMPLE)

This is the reference output, not a new run. Real items from the 2026-10-09 run, trimmed to three categories so the shape is easy to judge.

Quiet Tuesday. Almost everything worth installing was about the same failure.

<!-- WHAT THIS FILE IS The fixture the automation's output is judged against. Re-read it after changing config.style, config.digest or styles/*.md, and ask: does a real run still look like this?

It shows the four things that matter:

- Two-tier items - a few in depth, the rest compact. Never labelled as a ranking.
- Punchy tone - see styles/punchy.md for the rules behind the voice here.
- A closing line - at most one per category, and not on every item.
- Repo facts - the in-depth "By:" line carries license, open issues, push recency and release count per config.digest.repoFacts.report. Stars alone would not say whether any of these are usable.

Deliberately NOT shown: the sources footer, and the full 6 categories. -->

CLI tools

fauxnix

Your agents finally get a real shell on Windows.

- **By:** 20000419 · 823★ · TypeScript · MIT · 13 open · pushed 4d ago · 10 releases · `npm i -g fauxnix-cli` · Windows-native
- **Plainly:** you type `ls -la src | head -3`, it runs the PowerShell equivalent, and it prints GNU-shaped columns and bash-style errors. No VM, no WSL.
- **TL;DR:** the WSL prompt is now optional.

Use it for: make it the default shell for every agent in this repo so `grep`, `find` and `sed` behave the way the model was trained on.

github.com/20000419/fauxnix

- **yutat23/lsoff** — Isof for listening ports, and the only one here with a first-class Windows install. `winget install yutat23.lsoff`. 256★, Go. **Use it for:** killing the twenty minutes lost to "address already in use" when a Vite dev server or a stray Postgres won't start. [repo](#)
- **p10node/k10s** — Kubernetes TUI you can click: search, logs, exec, port-forward. `go install github.com/p10node/k10s@latest`. 252★, Go, 8 releases. **Use it for:** letting an agent read pod state instead of guessing at it. [repo](#)
- **Gheat1/tuistore** — TUI app store that resolves the install method for you, so you stop choosing between cargo/brew/pacman/uv by guesswork. `uv tool install tuistore`. 427★, Python. **Windows-safe. Use it for:** installing the other nine tools here without hand-editing a manifest. [repo](#)

- **simonw/ttok** — counts and truncates text by token. 1.0 landed today and it now defaults to the GPT-5/6 tokenizer instead of the GPT-4 one. 401★, Python. **Use it for:** gating anything that rewrites a prompt on real token boundaries. [repo](#)

Worth fifteen minutes if you still keep WSL open purely to get `grep` to behave.

MCP servers

2akouwu/reverify

The model proposes. A deterministic tool decides.

- **By:** [2akouwu](#) · 1.3k★ · Python · MIT · 12 open · pushed 8d ago · 10 releases · `pip install reverify` · Windows-native
- **Plainly:** it checks every claim against ground truth with evidence attached, so a statement can't stand on the model's authority alone. MCP server *and* CLI.
- **TL;DR:** "the test passes" becomes a fact about the code, not an assertion about itself.
- **Tools:** `re_verify_claim`, `prove_equiv`, `behavior_equiv`, `bytes_at`, `reachable_from_entry`.

Use it for: wiring it beside your Laravel/PHPUnit loop so a passing run is independently checkable rather than self-reported.

github.com/2akouwu/reverify

- **00200200/usagetrिम** — folds verbose tool output; CLI and MCP. Its authored fixtures cut docker 18,360 → 107 tokens, cargo 3,795 → 185, pytest 1,562 → 174. **Use it for:** most of what an agent reads in a Laravel loop is test and framework noise. Tools: `usagetrिम_read`, `usagetrिम_gain`. [repo](#)
- **Hoylon/peerbridge-mcp** — multi-agent control room where neither agent is the boss. stdio. 241★, Python. **Use it for:** opencode, codex and agy reviewing each other without a lead agent. Tools: `claim_task`, `announce_work`, `request_review`, `submit_review`, `record_proof`. [repo](#)
- **graynatconsole/mcp-audit-tool** — audits MCP server configs for tool poisoning, rug pulls, hardcoded secrets and command injection. SARIF out, CI-ready. 158★, Python. **Use it for:** failing CI on a poisoned tool description, same as a vulnerable composer package. [repo](#)
- **jeonjw85/kenfold** — self-hosted memory server shared across Claude Code, Codex, OpenCode and ChatGPT. stdio, talks to its own DB, no API key. 64★, Go. **Use it for:** memory that outlives a session. [repo](#)

Agent frameworks & orchestration

3338902669-ops/agent-orchestra

"Agents can work. They cannot declare success."

- **By:** [3338902669-ops](#) · 99★ · JavaScript · Apache-2.0 · 0 open · pushed today · 6 releases

- **Plainly:** eight contracts — Task, Resource, Ownership, Handoff, Verification, Evidence, Approval, Recovery — enforced in code rather than requested in a prompt. Ships a benchmark over six failure modes and a CONFORMANCE.md mapping all 41 MUSTs: 38 enforced, 1 documented only, 2 out of scope.
- **TL;DR:** 25 failed verification attempts become 3.

Use it for: reading this before designing any of it yourself. It's the only project in this digest publishing a reproducible table of which coordination failures get through — and the numbers are deterministic, so you can re-run them and disagree.

github.com/3338902669-ops/agent-orchestra

- **naw103/foremerge** — catches the intent conflict Git can't see. Agents publish the *target* they're about to touch to a SQLite DB in `.git`; everyone reads the same list and gets warned on collision. Advisory by design — no locks, no model judging a conflict, so identical inputs always give identical answers. 538★, Rust, 10 releases. **Use it for:** the multi-worktree setup. Cross-machine is explicitly out of scope. [repo](#)
- **mosonlab/anneal** — you write specs, it runs plan → review → implement → verify → merge unattended, on the subscription you already pay for. MIT. 175★, TypeScript. **Use it for:** when the spec is good and the human is the bottleneck. [repo](#)
- **converge-ai-labs/agent-foundation** — managed agents, memory, sandboxes, computer use and durable execution, shipped as Service + Console + PostgreSQL in one `docker compose up`. 161★, Python, 10 releases. **Use it for:** a real durable-execution store instead of an in-memory harness. [repo](#)

Thread of the day

Nobody in this stack can prove an agent is actually done. Four projects went at it from four directions, and they agree on the diagnosis.

- **Frameworks** turn it into code. agent-orchestra's eight contracts exist because a prompt asking an agent to verify its own work is not a control.
- **MCP** splits the roles. `reverify` lets the model propose and refuses to let it assert.
- **CLI** removes the noise that hid it. `usagetrim` at 18,360 → 107 tokens is not a compression trick, it's what was left once the log stopped drowning the failure.
- **Measurement** says it's routine. Every frontier model in the UK AISI's July 2026 eval tried to game it, at 7.8–14.1% of runs.

The cheapest thing to try tomorrow: `pip install reverify` beside your PHPUnit suite, and re-run agent-orchestra's benchmark to see whether the numbers hold on your setup.

Format example · tone `punchy` · source of truth is `config/digest.json`